

AQEVAL • EXECUTIVE SUMMARY

Evaluation session 2026-07-21 ? frozen AQUIPLEX build (AQUA engine)
NOT OFFICIAL ? NO BENCHMARK WAS EXECUTED; NO SCORES EXIST IN THIS REPORT

Benchmarks executed officially: 0 of 15. Scores reported: none.

Environment cannot perform model inference: egress policy returns 403
host_not_allowed for every provider AQUA routes to (OpenRouter, Groq,
Gemini, Together) and for gated-dataset hosts. Verified empirically
(evidence/egress-probes.txt). Provider keys exist in the frozen .env and
were deliberately not loaded (no reachable endpoint). Per the evaluation
rules, nothing was fabricated or estimated: every benchmark is reported
NOT EXECUTED with its exact reason.

VERIFIED THIS SESSION (NOT MODEL SCORES)

Framework selftest	PASS - official pass@k estimator, MATH equivalence, extractors; executor validated on official HumanEval canonical solutions (pass / wrong fails / loop times out)
Real runtime	frozen router booted sessionless, isolated data dir; real aquiplex adapter queried live engine 4x; failed only at provider stage; run stamped INVALID, no score
Datasets	GSM8K n=1319, HumanEval n=164, MBPP n=974 acquired from canonical sources, SHA-256 pinned; all others NOT ACQUIRED (reasons recorded)
Integrity fix	zero-scored runs now stamped INVALID in all report formats (only code change)

BENCHMARKS - ALL NOT EXECUTED

MMLU (reasoning)	dataset host also blocked (people.eecs.berkeley.edu)
MMLU-Pro (reasoning)	dataset on Hugging Face (blocked)
GPQA (reasoning)	dataset gated + HF blocked; anti-contamination zip password deliberately not automated
Humanity's Last Exam (reasoning)	dataset gated on HF (blocked); official scoring requires an external LLM judge (GPT-4o per HLE paper / o3-mini per Scale leaderboard) - no judge endpoint reachable
HumanEval (coding)	dataset ACQUIRED and verified; framework validated on official canonical solutions - only inference is blocked
MBPP (coding)	dataset ACQUIRED and verified - only inference is blocked
LiveCodeBench (coding)	current release_v6 identified; official harness + HF dataset required (blocked)
SWE-bench (coding)	official docker-based harness infrastructure unavailable in sandbox
GSM8K (math)	dataset ACQUIRED and verified - only inference is blocked
MATH-500 (math)	dataset on Hugging Face (blocked)
AIME (math)	dataset on Hugging Face (blocked)
MMMU (multimodal)	multimodal image-input adapter not yet wired + HF blocked
DocVQA (multimodal)	registration-gated dataset + multimodal adapter not yet wired
ChartQA (multimodal)	multimodal adapter not yet wired
BEIR (retrieval)	AQUA exposes no embeddings/retrieval endpoint - precondition missing regardless of environment

NEXT STEP

Run on a machine with open egress + the build's provider keys.

Exact commands: RUNBOOK.md. Only runs badged OFFICIAL are citable.

Scientific honesty over impressive numbers: this report contains no model performance figures because none were legitimately produced.