

AQEVAL • BENCHMARK RUN REPORT

gsm8k ? AQUA (Aquiplex Cognitive AI Operating System)
INVALID RUN ? ZERO ITEMS SCORED, NO BENCHMARK RESULT

exact_match_strict: 0.00%

stderr -0.00 • 95% CI [0.0, 0.0] • n=0

RUN RECORD

Run ID	20260721-144752_gsm8k_aquiplex
Date (UTC)	2026-07-21T14:48:07Z
Benchmark version	openai/grade-school-math @ master (test split, n=1319)
Adapter	aquiplex
Runtime	14.2 s
Seed	1234
Prompt template	gsm8k_cot_8shot gsm8k_cot_8shot/v1 sha256 3d55795c084e26f6
AQEval commit	134e27894155c10c67b5cef9b472b90d09f3283c
Platform commit	134e27894155c10c67b5cef9b472b90d09f3283c
Hardware	Linux-6.6.50-05083-g5a8e7e69b744-x86_64-with-glibc2.36 / 2 vCPU / 6.4 GB
Python	3.11.2
Errors	1 (100.00%)

DATASETS

test	datasets/cache/gsm8k/test.jsonl sha256 3730d312f6e3440559ace488
------	---

SUBSCORES

exact_match_flexible	0.00% (n=0)
----------------------	-------------

NOTE: RUN INVALID: 0 items were scored (100% errors). No benchmark result exists for this run ? see errors.log and records.jsonl.
NOTE: Non-official run (partial split, filtered subjects, or mock adapter). Not a reportable benchmark result.
Raw responses: records.jsonl • Full manifest: manifest.json • Generated by AQEval 1.0