

AQEVAL • BENCHMARK RUN REPORT

gsm8k ? AQUA (Aquiplex Cognitive AI Operating System)

SMOKE RUN ? PARTIAL / NON-DEFAULT, NOT REPORTABLE

exact_match_strict: 100.00%

stderr -0.00 • 95% CI [100.0, 100.0] • n=1

RUN RECORD

Run ID	20260721-151912_gsm8k_aquiplex
Date (UTC)	2026-07-21T15:19:25Z
Benchmark version	openai/grade-school-math @ master (test split, n=1319)
Adapter	aquiplex
Runtime	12.1 s
Seed	1234
Prompt template	gsm8k_cot_8shot gsm8k_cot_8shot/v1 sha256 3d55795c084e26f6
AQEval commit	134e27894155c10c67b5cef9b472b90d09f3283c
Platform commit	134e27894155c10c67b5cef9b472b90d09f3283c
Hardware	Linux-6.6.50-05083-g5a8e7e69b744-x86_64-with-glibc2.36 / 2 vCPU / 6.4 GB
Python	3.11.2
Errors	0 (0.00%)

DATASETS

test	datasets/cache/gsm8k/test.jsonl sha256 3730d312f6e3440559ace488
------	---

SUBSCORES

exact_match_flexible	100.00% (n=1)
----------------------	---------------

NOTE: Non-official run (partial split, filtered subjects, or mock adapter). Not a reportable benchmark result.

Raw responses: records.jsonl • Full manifest: manifest.json • Generated by AQEval 1.0