

AQEval • BENCHMARK RUN REPORT

humaneval ? mock-selftest

MOCK ? FRAMEWORK SELFTEST, NOT A MODEL RESULT

pass@1: 0.00%

stderr -0.00 • 95% CI [0.0, 0.0] • n=8

RUN RECORD

Run ID	20260722-075219_humaneval_mock_selftest
Date (UTC)	2026-07-22T07:52:20Z
Benchmark version	openai/human-eval HumanEval.jsonl.gz (n=164)
Adapter	mock
Runtime	0.1 s
Seed	1234
Prompt template	humaneval chat/v1 sha256 6a0f6c56f7ab5a73
AQEval commit	134e27894155c10c67b5cef9b472b90d09f3283c
Platform commit	134e27894155c10c67b5cef9b472b90d09f3283c
Hardware	Linux-6.6.50-05083-g5a8e7e69b744-x86_64-with-glibc2.36 / 2 vCPU / 6.4 GB
Python	3.11.2
Errors	0 (0.00%)

DATASETS

problems datasets/cache/humaneval/HumanEval.jsonl.gz sha256 b796127e635a67f93fb35c04

NOTE: Non-official run (partial split, filtered subjects, or mock adapter). Not a reportable benchmark result.

Raw responses: records.jsonl • Full manifest: manifest.json • Generated by AQEval 1.0